

bigdata_project_in_progress

April 12, 2024

1 Project: Big Data - Comparing ec2 Instance Performances

1.1 Short S3 Tutorial

- a) access s3 in your aws account
- b) upload data file
- c) access IAM in your aws account (put in aws search bar)
- d) look for ec2 default role (just familiarize yourself with the name)
- e) now go to your instances
- f) click on one of your instances' ID summaries
- g) click on the instant action drop down menu
- h) click on security
- i) select Modify IAM role
- j) add the default ec2 IAM role & then press update
- k) restart kernel in jupyter lab & run code below
** Ask Elisabeth or chatgpt for further clarification **

1.2 Connect Dask

```
[1]: from dask.distributed import Client

# Create a client and connect to the scheduler
client = Client('172.31.3.68:8786')

# Restart the client
client.restart()

# Perform an operation to check if the client is active
```

```
# For example, you can print the cluster information
print(client)
```

```
<Client: 'tcp://172.31.3.68:8786' processes=0 threads=0, memory=0 B>
```

1.3 Connect s3 data

```
[2]: import boto3
```

```
s3 = boto3.client('s3') # connect to s3
```

```
[4]: response = s3.list_objects_v2(Bucket='digit-dataset') # connect to s3 bucket
```

```
[5]: for obj in response['Contents']: # show bucket contents
      print(obj['Key'])
```

```
digits.csv
```

```
[8]: s3.download_file('digit-dataset', 'digits.csv', 'digits.csv') # should download
      ↪ into file successfully
```

```
[10]: import pandas as pd
       import numpy as np
```

```
[11]: df = pd.read_csv("digits.csv")
```

```
[14]: df.head(10)
```

```
[14]:  pixel_0_0  pixel_0_1  pixel_0_2  pixel_0_3  pixel_0_4  pixel_0_5  \
0         0.0         0.0         5.0         13.0         9.0         1.0
1         0.0         0.0         0.0         12.0         13.0         5.0
2         0.0         0.0         0.0          4.0         15.0        12.0
3         0.0         0.0         7.0         15.0         13.0         1.0
4         0.0         0.0         0.0          1.0         11.0         0.0
5         0.0         0.0        12.0        10.0          0.0         0.0
6         0.0         0.0         0.0        12.0        13.0         0.0
7         0.0         0.0         7.0          8.0        13.0        16.0
8         0.0         0.0         9.0        14.0          8.0         1.0
9         0.0         0.0        11.0        12.0          0.0         0.0

      pixel_0_6  pixel_0_7  pixel_1_0  pixel_1_1  ...  pixel_6_7  pixel_7_0  \
0         0.0         0.0         0.0         0.0  ...         0.0         0.0
1         0.0         0.0         0.0         0.0  ...         0.0         0.0
2         0.0         0.0         0.0         0.0  ...         0.0         0.0
3         0.0         0.0         0.0         8.0  ...         0.0         0.0
4         0.0         0.0         0.0         0.0  ...         0.0         0.0
5         0.0         0.0         0.0         0.0  ...         0.0         0.0
6         0.0         0.0         0.0         0.0  ...         0.0         0.0
```

7	15.0	1.0	0.0	0.0	...	0.0	0.0
8	0.0	0.0	0.0	0.0	...	0.0	0.0
9	0.0	0.0	0.0	2.0	...	0.0	0.0

	pixel_7_1	pixel_7_2	pixel_7_3	pixel_7_4	pixel_7_5	pixel_7_6	\
0	0.0	6.0	13.0	10.0	0.0	0.0	
1	0.0	0.0	11.0	16.0	10.0	0.0	
2	0.0	0.0	3.0	11.0	16.0	9.0	
3	0.0	7.0	13.0	13.0	9.0	0.0	
4	0.0	0.0	2.0	16.0	4.0	0.0	
5	0.0	9.0	16.0	16.0	10.0	0.0	
6	0.0	1.0	9.0	15.0	11.0	3.0	
7	0.0	13.0	5.0	0.0	0.0	0.0	
8	0.0	11.0	16.0	15.0	11.0	1.0	
9	0.0	9.0	12.0	13.0	3.0	0.0	

	pixel_7_7	target
0	0.0	0
1	0.0	1
2	0.0	2
3	0.0	3
4	0.0	4
5	0.0	5
6	0.0	6
7	0.0	7
8	0.0	8
9	0.0	9

[10 rows x 65 columns]

```
[34]: import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap

import sklearn
from sklearn.inspection import DecisionBoundaryDisplay
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier, \
    NeighborhoodComponentsAnalysis
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
import seaborn as sns
from sklearn.metrics import classification_report
```

```
[28]: X_train, X_test, y_train, y_test = train_test_split(df, df.target, test_size=0.
    ↪2, random_state=42)
```

```
[29]: knn = KNeighborsClassifier(n_neighbors=5)
```

```
knn.fit(X_train, y_train)
```

```
KNeighborsClassifier()
```

```
y_pred = knn.predict(X_test)
```

```
classification_report(y_test, y_pred)
```

	precision	recall	f1-score	support\n\n	0		
1.00	1.00	1.00	33\n	1	1.00	1.00	1.00
28\n	2	1.00	1.00	1.00	33\n	3	
1.00	1.00	1.00	34\n	4	0.98	1.00	0.99
46\n	5	0.98	0.96	0.97	47\n	6	
0.97	1.00	0.99	35\n	7	1.00	0.97	0.99
34\n	8	1.00	1.00	1.00	30\n	9	
0.95	0.95	0.95	40\n\n	accuracy			
0.99	360\n	macro avg	0.99	0.99	0.99	360\n	nweighted
avg	0.99	0.99	0.99	360\n'			

```
k_val = range(1,11)

# Initialize lists to store accuracy scores
train_accuracy = []
test_accuracy = []

# Iterate over each value of k
for k in k_val:


    knn = KNeighborsClassifier(n_neighbors=k)

    knn.fit(X_train, y_train)

    train_accuracy.append(knn.score(X_train, y_train))

    test_accuracy.append(knn.score(X_test, y_test))

# Plot the accuracy scores
plt.plot(k_val, train_accuracy, label='Train Accuracy')
plt.plot(k_val, test_accuracy, label='Test Accuracy')
plt.xlabel('Number of Neighbors (k)')
plt.ylabel('Accuracy')
plt.title('Accuracy vs. Number of Neighbors')
plt.legend()
plt.grid(True)
plt.show()
```



```
[53]: for k_val, train_acc, test_acc in zip(k_val, train_accuracy, test_accuracy):
        print(f'k={k_val}: Train Accuracy={train_acc:.2f}, Test Accuracy={test_acc:.2f}')
```

```
k=1: Train Accuracy=1.00, Test Accuracy=0.98
k=2: Train Accuracy=0.99, Test Accuracy=0.98
k=3: Train Accuracy=0.99, Test Accuracy=0.99
k=4: Train Accuracy=0.99, Test Accuracy=0.98
k=5: Train Accuracy=0.99, Test Accuracy=0.99
k=6: Train Accuracy=0.99, Test Accuracy=0.99
k=7: Train Accuracy=0.99, Test Accuracy=0.99
k=8: Train Accuracy=0.99, Test Accuracy=0.99
k=9: Train Accuracy=0.99, Test Accuracy=0.99
k=10: Train Accuracy=0.99, Test Accuracy=0.98
```

```
[55]: from sklearn.model_selection import train_test_split
        from sklearn.neural_network import MLPClassifier
```

```
[56]: mlp = MLPClassifier(hidden_layer_sizes=(2,), activation='logistic',
        max_iter=1000)
```

```
[57]: # Train the XOR classifier
mlp.fit(X_train, y_train)
```

```
/home/ubuntu/.local/lib/python3.10/site-
packages/sklearn/neural_network/_multilayer_perceptron.py:691:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (1000) reached and
the optimization hasn't converged yet.
warnings.warn(
```

```
[57]: MLPClassifier(activation='logistic', hidden_layer_sizes=(2,), max_iter=1000)
```

```
[58]: # Predict the labels for the test data
y_pred = mlp.predict(X_test)
```

```
[59]: # Evaluate the classifier's performance
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.33	0.03	0.06	33
1	0.46	0.93	0.61	28
2	0.43	0.88	0.57	33
3	0.77	0.29	0.43	34
4	0.86	0.26	0.40	46
5	0.00	0.00	0.00	47
6	0.34	0.91	0.50	35
7	0.57	0.24	0.33	34
8	0.14	0.07	0.09	30
9	0.45	0.93	0.60	40
accuracy			0.44	360
macro avg	0.43	0.45	0.36	360
weighted avg	0.44	0.44	0.35	360


```
/home/ubuntu/.local/lib/python3.10/site-
packages/sklearn/metrics/_classification.py:1509: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/ubuntu/.local/lib/python3.10/site-
packages/sklearn/metrics/_classification.py:1509: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/ubuntu/.local/lib/python3.10/site-
packages/sklearn/metrics/_classification.py:1509: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
[60]: from sklearn.metrics import confusion_matrix
```

```
[61]: # Calculate confusion matrix for XOR model
conf_matrix_xor = confusion_matrix(y_test, y_pred)

# Plot confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix_xor, annot=True, cmap='Blues', fmt='g', cbar=False)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix for XOR Model')
plt.show()
```



```
[63]: import tensorflow as tf
from tensorflow.keras import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
```

2024-04-12 08:34:23.971769: I tensorflow/core/platform/cpu_feature_guard.cc:210]

This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: AVX2 AVX512F FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
2024-04-12 08:34:25.927707: W
tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning: Could not find TensorRT

```
[65]: # Reshape the digit images for CNN input
X_train_cnn = X_train.values.reshape(-1, 8, 8, 1)
X_test_cnn = X_test.values.reshape(-1, 8, 8, 1)
```

```
-----
ValueError                                Traceback (most recent call last)
Cell In[65], line 2
      1 # Reshape the digit images for CNN input
----> 2 X_train_cnn = X_train.values.reshape(-1, 8, 8, 1)
      3 X_test_cnn = X_test.values.reshape(-1, 8, 8, 1)

ValueError: cannot reshape array of size 93405 into shape (8,8,1)
```

```
[ ]: # Define the CNN model
model = Sequential([
    Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(8, 8, 1)),
    MaxPooling2D(pool_size=(2, 2)),
    Flatten(),
    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])
```

```
[ ]: # Compile the model
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

```
[ ]: # Train the model
model.fit(X_train_cnn, y_train, epochs=10, batch_size=32,
          validation_data=(X_test_cnn, y_test))
```

```
[ ]: # Evaluate the model on test data
test_loss, test_acc = model.evaluate(X_test_cnn, y_test)
print(f'Test accuracy: {test_acc}')
```