

Install and Import Neccessary Packages

```
In [ ]: # !pip install dask-expr
        # !pip install s3fs
        # !pip install boto3
        # %pip install matplotlib
        # !pip3 install scikit-learn
        # %pip install dask
        # %pip install seaborn
        # %pip install tensorflow
```

```
In [18]: import dask
        from dask.distributed import Client
        import dask.dataframe as dd
        import boto3
        import matplotlib.pyplot as plt
        from matplotlib.colors import ListedColormap
        import sklearn
        from sklearn.inspection import DecisionBoundaryDisplay
        from sklearn.model_selection import train_test_split
        from sklearn.neighbors import KNeighborsClassifier, NeighborhoodComponentsAr
        from sklearn.pipeline import Pipeline
        from sklearn.preprocessing import StandardScaler
        import seaborn as sns
        from sklearn.metrics import classification_report
        import time
        import tensorflow as tf
        from tensorflow.keras import Sequential
        from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
        import pandas as pd
        import numpy as np
```

Connect to Client

```
In [35]: client = Client('172.31.10.249:8786')

        # Restart the client
        client.restart()

        print(client)
```

<Client: 'tcp://172.31.10.249:8786' processes=1 threads=1, memory=3.81 GiB>

```
In [36]: workers = client.scheduler_info()['workers']
        print("Number of workers:", len(workers))

        # Optionally, print details about each worker
        for worker, details in workers.items():
            print(f"Worker {worker}:")
            print("  - Host:", details['host'])
```

```
print(" - Number of threads:", details['nthreads'])
print(" - Memory limit:", details['memory_limit'], "bytes")
```

```
Number of workers: 1
Worker tcp://172.31.10.174:43531:
  - Host: 172.31.10.174
  - Number of threads: 1
  - Memory limit: 4095528960 bytes
```

```
In [37]: s3 = boto3.client('s3') # connect to s3
```

```
In [38]: response = s3.list_objects_v2(Bucket='digit-dataset') # connect to s3 bucket
```

```
In [39]: for obj in response['Contents']: # show bucket contents
          print(obj['Key'])
```

```
digits.csv
multi-digit.csv
```

```
In [40]: s3_path = 's3://digit-dataset/digits.csv'
```

```
In [41]: # !pip install pickleshare
```

```
In [42]: import os
          os.chdir('/home/ubuntu')
```

```
In [43]: cd '/home/ubuntu'
```

```
/home/ubuntu
```

```
In [44]: ls
```

```
Big_Data_Project.ipynb  digits.csv
```

```
In [45]: df = dd.read_csv(s3_path).sample(frac=0.5) # Large dataset, use only part
```

```
In [46]: df
```

```
Out[46]: Dask DataFrame Structure:
```

```
           pixel_0_0  pixel_0_1  pixel_0_2  pixel_0_3  pixel_0_4  pixel_0_5  pixel_0_6
npartitions=1
```

```
           float64      float64      float64      float64      float64      float64      float64
```

```
           ...           ...           ...           ...           ...           ...           ...
```

```
Dask Name: sample, 2 expressions
```

```
In [47]: sampled_df.info()
```

```
<class 'dask_expr.DataFrame'>
Columns: 65 entries, pixel_0_0 to target
dtypes: float64(64), int64(1)
```

Split Data

```
In [48]: import dask.array as da
import dask.dataframe as dd
from sklearn.datasets import load_digits
```

```
In [49]: # !pip install dask_ml
```

```
In [50]: X = df.drop(columns=['target']).compute() # Features
y = df['target'].compute() # Target variable
```

```
In [52]: X_dask = da.from_array(X.values.reshape(-1, 8, 8, 1), chunks=(1000, 8, 8, 1))
y_dask = da.from_array(y.values, chunks=(1000,))
```

```
In [58]: X_flattened = X_dask.reshape(X_dask.shape[0], -1).compute()
```

```
In [59]: X_train, X_test, y_train, y_test = train_test_split(X_flattened, y_dask, tes
```

KNN Model

```
In [60]: %%time
knn = KNeighborsClassifier(n_neighbors=5)
```

CPU times: user 13 µs, sys: 1 µs, total: 14 µs
Wall time: 17.2 µs

```
In [62]: %%time
# fit KNN model
knn.fit(X_train, y_train)
```

CPU times: user 11.7 ms, sys: 0 ns, total: 11.7 ms
Wall time: 34.6 ms

```
Out[62]: KNeighborsClassifier
KNeighborsClassifier()
```

```
In [63]: %%time
y_pred = knn.predict(X_test)
```

CPU times: user 31.4 ms, sys: 13 ms, total: 44.4 ms
Wall time: 60 ms

```
In [64]: %%time
classification_report(y_test, y_pred)
```

CPU times: user 44.2 ms, sys: 7.45 ms, total: 51.6 ms
Wall time: 136 ms

```
Out[64]: '
           precision    recall  f1-score   support\n\n
    0.96      1.00      1.00      1.00      12\n
    0.96      0.96      0.95      0.96      1\n
    3      1.00      0.94      0.97      16\n
    00      1.00      1.00      1.00      4\n
    16\n      6      1.00      1.00      1.00      28\n
    0.95      1.00      0.97      1.00      8\n
    0.93      15\n      9      1.00      0.88      0.93      16\n\n
accuracy                    0.97      180\n
0.97      0.97      180\nweighted avg      0.97      0.97      0.97
180\n'
```

```
In [66]: %%time
k_val = range(1,11)

# Initialize lists to store accuracy scores
train_accuracy = []
test_accuracy = []
train_loss = []
test_loss = []

# Iterate over each value of k
for k in k_val:

    knn = KNeighborsClassifier(n_neighbors=k)

    knn.fit(X_train, y_train)

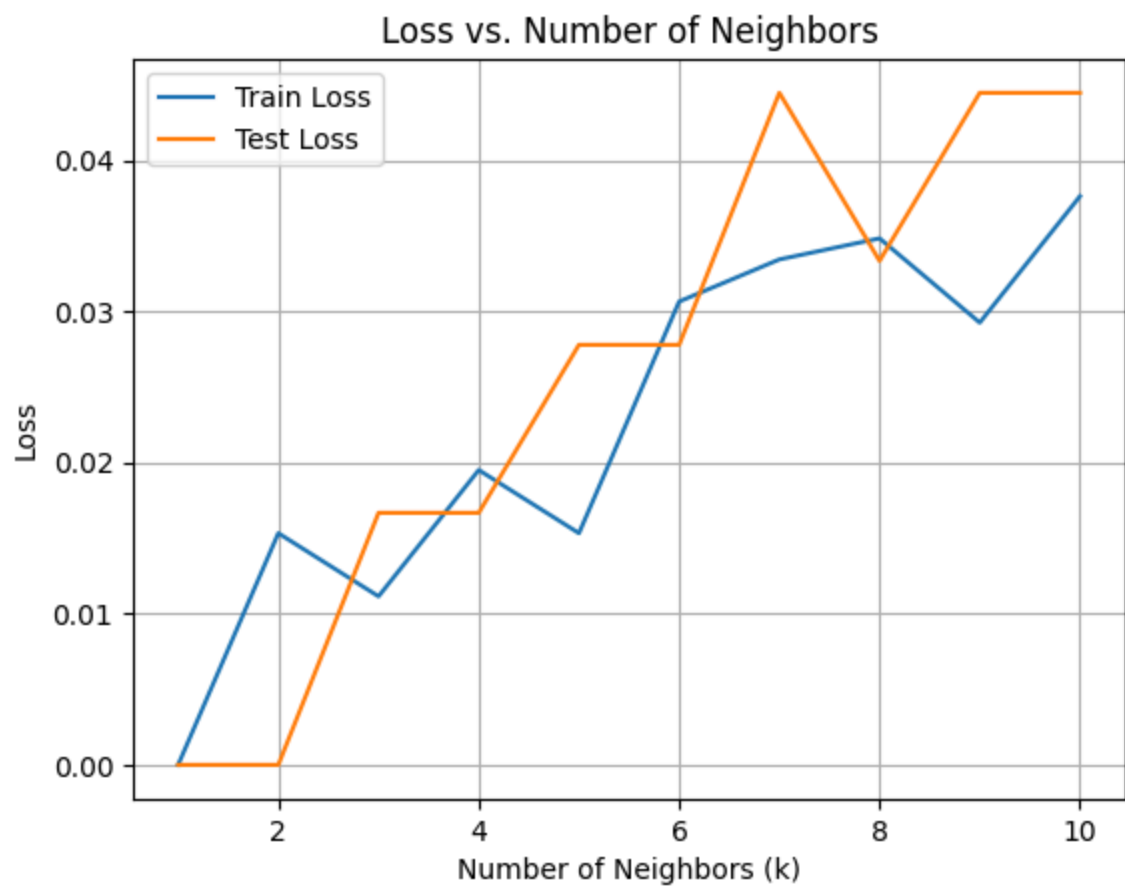

    train_accuracy.append(knn.score(X_train, y_train))
    test_accuracy.append(knn.score(X_test, y_test))

    train_pred = knn.predict(X_train)
    test_pred = knn.predict(X_test)

    train_loss.append(np.mean(train_pred != y_train)) # Classification error
    test_loss.append(np.mean(test_pred != y_test)) # Classification error

# Plot the accuracy scores
plt.plot(k_val, train_accuracy, label='Train Accuracy')
plt.plot(k_val, test_accuracy, label='Test Accuracy')
plt.xlabel('Number of Neighbors (k)')
plt.ylabel('Accuracy')
plt.title('Accuracy vs. Number of Neighbors')
plt.legend()
plt.grid(True)

plt.show()
plt.plot(k_val, train_loss, label='Train Loss')
plt.plot(k_val, test_loss, label='Test Loss')
plt.xlabel('Number of Neighbors (k)')
plt.ylabel('Loss')
plt.title('Loss vs. Number of Neighbors')
plt.legend()
plt.grid(True)
plt.show()
```



CPU times: user 1.93 s, sys: 146 ms, total: 2.07 s
Wall time: 4.65 s

CNN Model

```
In [67]: X_train_cnn = X_train.reshape(-1, 8, 8, 1).astype(np.float32)
X_test_cnn = X_test.reshape(-1, 8, 8, 1).astype(np.float32)
```

```
In [68]: %%time
# Define the CNN model
model = Sequential([
    Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(8, 8, 1))
    MaxPooling2D(pool_size=(2, 2)),
    Flatten(),
    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])
```

CPU times: user 68.6 ms, sys: 0 ns, total: 68.6 ms

Wall time: 78.8 ms

```
/home/ubuntu/.local/lib/python3.10/site-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
In [69]: %%time
# Compile the model
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

CPU times: user 7.86 ms, sys: 4.45 ms, total: 12.3 ms

Wall time: 10.8 ms

```
In [70]: %%time
# Train the model
model.fit(X_train_cnn, y_train, epochs=10, batch_size=10, validation_data=(X_test_cnn, y_test))
```

```

Epoch 1/10
72/72 ————— 1s 5ms/step - accuracy: 0.4372 - loss: 2.2326 - v
al_accuracy: 0.8333 - val_loss: 0.5612
Epoch 2/10
72/72 ————— 0s 2ms/step - accuracy: 0.9391 - loss: 0.2913 - v
al_accuracy: 0.9389 - val_loss: 0.2519
Epoch 3/10
72/72 ————— 0s 2ms/step - accuracy: 0.9599 - loss: 0.1556 - v
al_accuracy: 0.9389 - val_loss: 0.2143
Epoch 4/10
72/72 ————— 0s 3ms/step - accuracy: 0.9692 - loss: 0.1108 - v
al_accuracy: 0.9611 - val_loss: 0.1346
Epoch 5/10
72/72 ————— 0s 3ms/step - accuracy: 0.9893 - loss: 0.0681 - v
al_accuracy: 0.9556 - val_loss: 0.1061
Epoch 6/10
72/72 ————— 0s 3ms/step - accuracy: 0.9960 - loss: 0.0413 - v
al_accuracy: 0.9611 - val_loss: 0.1160
Epoch 7/10
72/72 ————— 0s 3ms/step - accuracy: 0.9997 - loss: 0.0274 - v
al_accuracy: 0.9722 - val_loss: 0.1301
Epoch 8/10
72/72 ————— 0s 4ms/step - accuracy: 0.9992 - loss: 0.0254 - v
al_accuracy: 0.9667 - val_loss: 0.0899
Epoch 9/10
72/72 ————— 0s 4ms/step - accuracy: 1.0000 - loss: 0.0137 - v
al_accuracy: 0.9944 - val_loss: 0.0842
Epoch 10/10
72/72 ————— 0s 3ms/step - accuracy: 1.0000 - loss: 0.0146 - v
al_accuracy: 0.9722 - val_loss: 0.0775
CPU times: user 4.08 s, sys: 145 ms, total: 4.23 s
Wall time: 4.13 s

```

Out[70]: <keras.src.callbacks.history.History at 0x7df7f4cc4af0>

```

In [71]: # Evaluate the model on test data
test_loss, test_acc = model.evaluate(X_test_cnn, y_test)
print(f'Test accuracy: {test_acc}, Test loss: {test_loss}')

```

```

6/6 ————— 0s 3ms/step - accuracy: 0.9686 - loss: 0.0833
Test accuracy: 0.9722222089767456, Test loss: 0.07749420404434204

```

```

In [72]: %%time
# Make predictions on validation data
predictions = model.predict(X_test_cnn)

```

```

6/6 ————— 0s 11ms/step
CPU times: user 198 ms, sys: 8.98 ms, total: 207 ms
Wall time: 228 ms

```

```

In [73]: # Plot images in a grid
num_images_to_plot = 5
num_cols = 5 # Number of columns in the grid
num_rows = (num_images_to_plot - 1) // num_cols + 1 # Calculate number of r

# Adjust figsize as needed
fig, axes = plt.subplots(num_rows, num_cols, figsize=(12, 12))

```

```

for i, ax in enumerate(axes.flat):
    if i < num_images_to_plot:
        # Plot the original image
        ax.imshow(X_test_cnn[i].reshape(8, 8), cmap='gray')
        ax.axis('off')

        # Get the predicted label for the current image
        predicted_label = np.argmax(predictions[i])


        # Get the actual label for the current image
        actual_label = y_test.compute()[i]

        # Set the title with predicted and actual labels
        ax.set_title(f"Predicted: {predicted_label}, Actual: {actual_label}")
    else:
        ax.axis('off') # Turn off empty subplots

plt.tight_layout() # Adjust spacing between subplots
plt.show()

```

2024-05-02 01:22:11,491 - distributed.client - WARNING - Couldn't gather 1 keys, rescheduling (('getitem-538764e8ef34b4752781504096b42cb0', 0),)



In []: